

OTG Legal Box — Product Brief

For managing partners, IT heads, and procurement | April 2026

At a Glance

Product	OTG Legal Box — a fully local AI assistant for Singapore law firms
Core capability	10 built-in legal tools powered by Gemma 4, running entirely on-premises
Data residency	All processing occurs on the firm's own hardware; no data leaves the device
Privacy	Automatic PII Shield anonymises names, NRICs, UENs, phones, emails, and bank accounts before the AI processes any text
Compliance	Supports PDPA accountability obligations; no third-party data processing agreement required
Deployment	Desktop (per device) or Server mode (shared Mac Mini/Mac Studio for 5–15 users)
AI model	Google DeepMind Gemma 4 — three tiers (E4B, 26B, 31B)
Connectivity	No internet required after initial setup; offline-first by design
Updates	USB delivery available for fully air-gapped environments

The Problem: Cloud AI and Legal Confidentiality

Large language models offer significant productivity gains for legal work — summarising judgments, reviewing contracts, and drafting correspondence in minutes rather than hours. However, mainstream AI tools are cloud-based: text is transmitted to remote servers, processed by third parties, and may be stored or used for training.

For Singapore law firms, this creates three interlocking risks:

- **Legal professional privilege** — uploading privileged documents to a cloud AI raises genuine questions about waiver of privilege through third-party disclosure.
- **PDPA obligations** — transferring personal data to an overseas data intermediary requires comparable protection standards, which are difficult to verify and maintain with large cloud AI providers.

- **Data residency** — clients (particularly financial institutions, government-linked entities, and MNCs) increasingly require that their data not leave Singapore.

The result: firms that use cloud AI expose themselves to risk; firms that ban AI entirely forgo a significant competitive advantage.

The Solution: OTG Legal Box

OTG Legal Box resolves this dilemma by running the entire AI stack on the firm's own hardware. The Gemma 4 model is downloaded once and stored locally. All inference happens on-device. No query, document, or response is transmitted over the internet.

10 Built-In Legal Tools

Tool	Description
Legal Chat	Conversational AI for legal research, statutory interpretation, and procedural questions
Summarize	Structured summaries of contracts, judgments, affidavits, and correspondence
Upload Cases	Index PDF/Word documents into a searchable local case library
Search Cases	Semantic search across indexed documents using natural language
Contract Review	Risk analysis identifying unusual clauses and missing protections
Redlining	Tracked-changes comparison with plain-English summary
Bundle Creator	Automated court bundle assembly — bookmarked, paginated PDFs
Chronology	Multi-document timeline extraction into an editable chronology
Drafting	First-draft generation for letters, billing narratives, and pleadings
Prompt Library	Reusable AI instruction templates to standardise firm workflows

PII Shield — Automatic Anonymisation

Before any text reaches the AI model, the PII Shield replaces sensitive data with neutral placeholders:

- Names → [PERSON_1] , NRICs → [NRIC_1] , Phone numbers → [PHONE_1]
- Emails → [EMAIL_1] , UENs → [UEN_1] , Bank accounts → [BANK_ACCT_1]

After the AI responds, placeholders are restored to original values. The model never processes real client identifiers.

Audit Log

Every action is recorded in a tamper-evident local audit log — timestamps, user identity, tool used, PII detection status, and network status (`local_only`). Exportable as CSV or PDF for compliance reviews.

Data Security & PDPA Compliance

- **No outbound connections** — no telemetry, analytics, licence checks, or API calls during operation
- **No third-party processor** — the firm is both data controller and processor for AI work
- **Section 26 (Transfer limitation)** — not engaged, as no data leaves the premises
- **Section 11 (Accountability)** — the audit log provides contemporaneous, exportable evidence of all AI processing, PII handling, and network status
- **No data processing agreement required** — eliminates the operational complexity of cloud AI compliance

System Requirements

All three tiers deliver the same ten tools. The difference is response speed, document capacity, and context window.

	Minimum	Recommended	Professional
Model	Gemma 4 E4B	Gemma 4 26B	Gemma 4 31B
RAM	8 GB	16 GB	32 GB
Storage	20 GB free	80 GB free	106 GB free
Hardware	Any laptop (2017+)	Modern laptop / Mac Mini M4	Mac Studio M4 Max/ Ultra
Context window	128K tokens (~90K words)	256K tokens (~180K words)	256K tokens (~180K words)
OS	macOS 14+ / Windows 10+	macOS 14+ / Windows 10+	macOS 14+
GPU	Not required (CPU mode)	Apple Silicon or NVIDIA RTX 3060+	Apple Silicon

Apple Silicon Macs (M1/M2/M3/M4) deliver the best performance due to unified memory architecture. **Windows PCs** with NVIDIA GPUs are fully supported via CUDA. Used hardware works — a 16 GB laptop with the E4B model runs all tools comfortably.

Deployment Options

	Desktop Mode	Server Mode
Installation	<code>.dmg</code> (Mac) or <code>.exe</code> (Windows)	<code>.pkg</code> on Mac Mini or Mac Studio
Users	1 per device	5–15 on office Wi-Fi
Hardware	Lawyer's own laptop or desktop	Dedicated Mac in office
Access	Standalone desktop app	Browser-based on office Wi-Fi (no install on individual devices)
User management	Per-device profiles	Central user management
Shared resources	Individual only	Shared case library; individual chat histories

Desktop Mode — best for sole practitioners, small firms (< 5 lawyers), or lawyers needing AI when working remotely.

Server Mode — best for firms with multiple lawyers sharing a document library, or where centralised IT administration is preferred.

Gemma 4 Model Tiers

E4B (Minimum) — 8 GB RAM, ~5 GB storage. Runs on older laptops. Suited to straightforward tasks: document summaries, letter drafting, common research questions. Best entry point for sole practitioners.

26B (Recommended) — 16 GB RAM, ~20 GB storage. Substantially improved reasoning and output quality. Handles complex documents comfortably. Best fit for most firms and all Server Mode deployments.

31B (Professional) — 32 GB RAM, ~25 GB storage. Highest output quality for complex multi-party pleadings, large transaction documents, and heavy discovery sets. Recommended hardware: Mac Studio. Supports 5–15 concurrent users in Server Mode.

Proof of Concept Offer

For the first five firms, OTG is offering a streamlined POC engagement:

Installation & setup	2 man-days on-site
Support	Up to 10 visits over first 3 months
POC cost	SGD \$2,000 (hardware not included)
Hardware	Firm purchases own hardware — can start with a Mac Mini; recommended to evaluate upcoming Apple Silicon releases for optimal LLM performance
Scope	Firm selects tools and workflows to trial

This POC allows the firm to evaluate OTG Legal Box in a live environment with real workflows before committing to a full deployment.

Beyond POC

Rate	SGD \$250/hour (SGD \$2,000/man-day)
Hardware	Purchased separately by the firm
Includes	Installation, configuration, onboarding, and ongoing support

Contact

For enquiries, demonstrations, or to reserve a POC slot, contact On The Ground AI.

[Contact details to be provided by OTG]